

RAY TRACING IN ONE WEEKEND: THE BOOK SERIES

Notes by Yuxing Shi

<https://raytracing.github.io/>

Ray Tracing in One Weekend

<https://raytracing.github.io/books/RayTracingInOneWeekend.html>

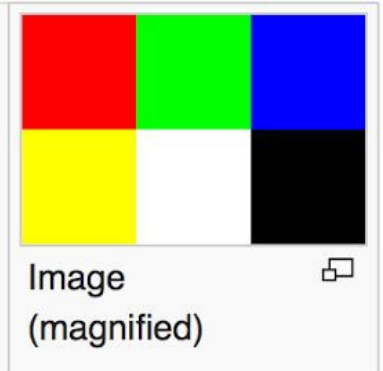
Output an Image

- The PPM Image Format

PPM example [\[edit\]](#)

This is an example of a color RGB image stored in PPM format. There is a newline character at the end of each line.

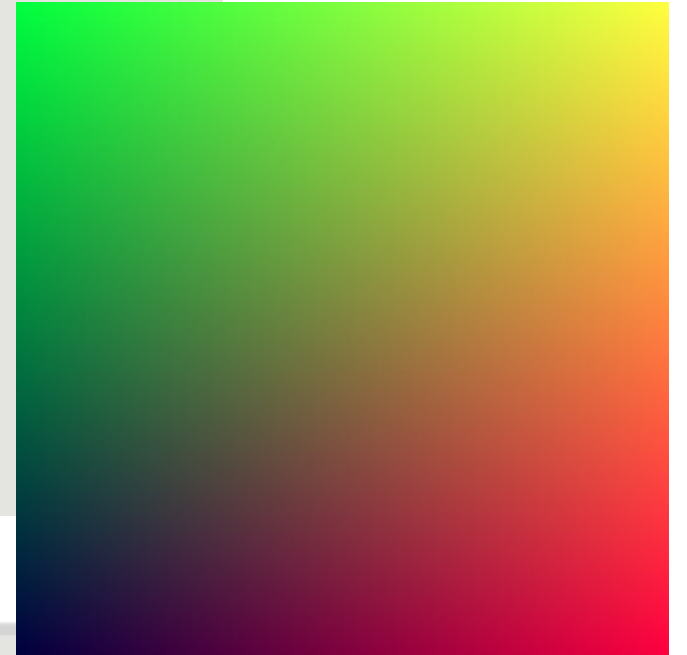
```
P3
# The P3 means colors are in ASCII, then 3 columns and 2 rows,
# then 255 for max color, then RGB triplets
3 2
255
255 0 0 0 255 0 0 0 255
255 255 0 255 255 255 0 0 0
```



```
// Render
```

```
std::cout << "P3\n" << image_width << ' ' << image_height << "\n255\n";
```

```
for (int j = image_height-1; j >= 0; --j) {  
    for (int i = 0; i < image_width; ++i) {  
        auto r = double(i) / (image_width-1);  
        auto g = double(j) / (image_height-1);  
        auto b = 0.25;  
  
        int ir = static_cast<int>(255.999 * r);  
        int ig = static_cast<int>(255.999 * g);  
        int ib = static_cast<int>(255.999 * b);  
  
        std::cout << ir << ' ' << ig << ' ' << ib << '\n';  
    }  
}
```



- Progress Indicator

```
std::cerr << "\nDone.\n";
```

```
build/inOneWeekend > image.ppm
```

Rays, a simple camera, and bg.

4.1. The ray Class

The one thing that all ray tracers have is a ray class and a computation of what color is seen along a ray. Let's think of a ray as a function $\mathbf{P}(t) = \mathbf{A} + t\mathbf{b}$. Here \mathbf{P} is a 3D position along a line in 3D. \mathbf{A} is the ray origin and \mathbf{b} is the ray direction. The ray parameter t is a real number (double in the code). Plug in a different t and $\mathbf{P}(t)$ moves the point along the ray. Add in negative t values and you can go anywhere on the 3D line. For positive t , you get only the parts in front of \mathbf{A} , and this is what is often called a half-line or ray.

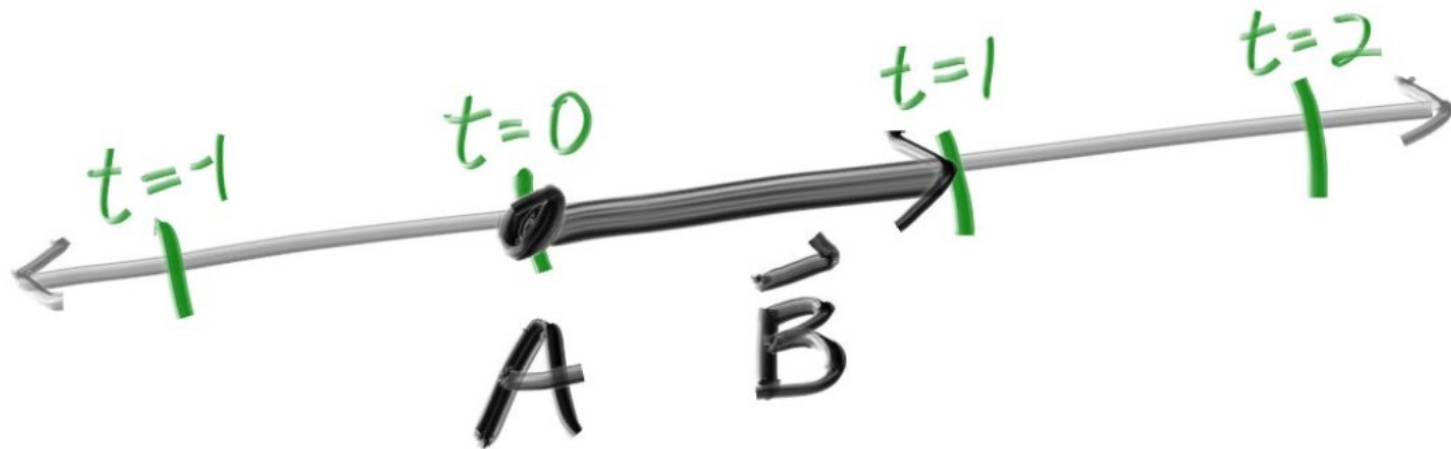
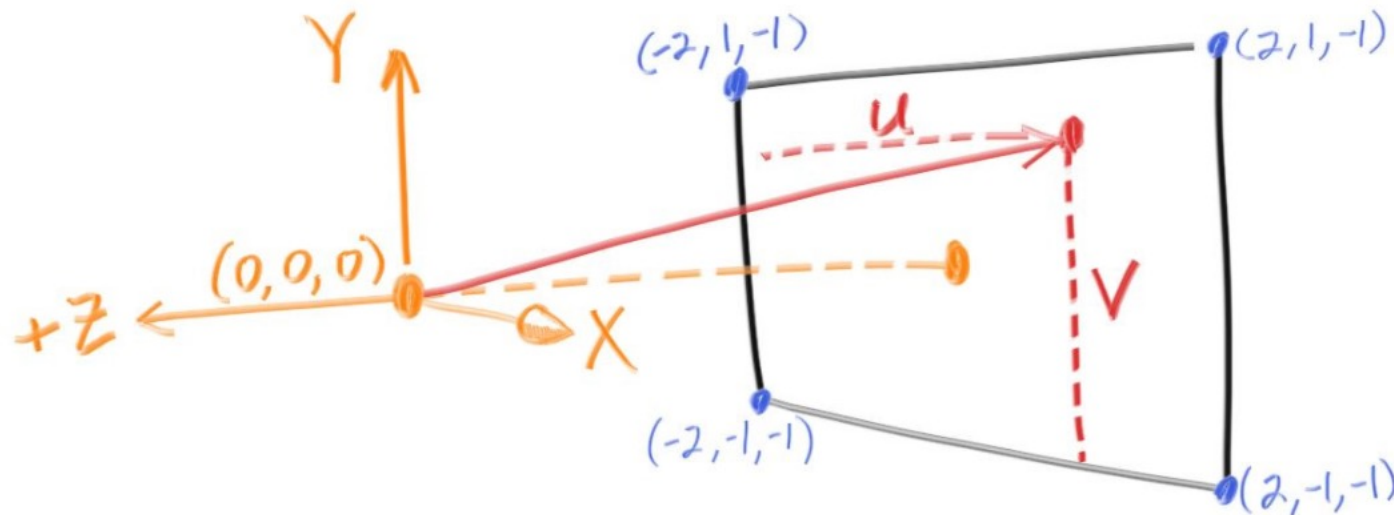


Figure 2: Linear interpolation

Sending Rays Into the Scene

- (1) calculate the ray from the eye to the pixel,
- (2) determine which objects the ray intersects, and
- (3) compute a color for that intersection point.



Focal length = 1 unit (orange)
!= focus distance

Figure 3: Camera geometry

Codes

- Viewpoint
 - through which to pass our scene rays
 - aspect ratio should be the same as our rendered image
 - Rectangle height of 2 at last slide.
 - Focal length: projection point to projection plane.

```
color ray_color(const ray& r) {  
    vec3 unit_direction = unit_vector(r.direction());  
    auto t = 0.5*(unit_direction.y() + 1.0);  
    return (1.0-t)*color(1.0, 1.0, 1.0) + t*color(0.5, 0.7, 1.0);  
}
```

- linearly blends white and blue depending on the height of the y coordinate after scaling the ray direction to unit length $(-1 \sim 1)$
- And then scale $0 \leq t \leq 1.0$

Codes

- Image + Camera (+ Render)

```
// Image
const auto aspect_ratio = 16.0 / 9.0;
const int image_width = 400;
const int image_height = static_cast<int>(image_width / aspect_ratio);

// Camera

auto viewport_height = 2.0;
auto viewport_width = aspect_ratio * viewport_height;
auto focal_length = 1.0;

auto origin = point3(0, 0, 0);
auto horizontal = vec3(viewport_width, 0, 0);
auto vertical = vec3(0, viewport_height, 0);
auto lower_left_corner = origin - horizontal/2 - vertical/2 - vec3(0, 0, focal_length);
```


Codes

- (Image + Camera +) Render

```
// Render
```

```
std::cout << "P3\n" << image_width << " " << image_height << "\n255\n";

for (int j = image_height-1; j >= 0; --j) {
    std::cerr << "\rScanlines remaining: " << j << ' ' << std::flush;
    for (int i = 0; i < image_width; ++i) {
        auto u = double(i) / (image_width-1);
        auto v = double(j) / (image_height-1);
        ray r(origin, lower_left_corner + u*horizontal + v*vertical - origin);
        color pixel_color = ray_color(r);
        write_color(std::cout, pixel_color);
    }
}
```



Image 2: A blue-to-white gradient depending on ray Y coordinate

Adding a Sphere

- Sphere: $(x-Cx)^2+(y-Cy)^2+(z-Cz)^2=r^2$
- Vector form: $(\mathbf{P}-\mathbf{C})\cdot(\mathbf{P}-\mathbf{C})=r^2$
- Ray hits sphere: $(\mathbf{P}(t)-\mathbf{C})\cdot(\mathbf{P}(t)-\mathbf{C})=r^2$
Where: $\mathbf{P}(t)=\mathbf{A}+t\mathbf{b}$
- $\Rightarrow t^2\mathbf{b}\cdot\mathbf{b}+2t\mathbf{b}\cdot(\mathbf{A}-\mathbf{C})+(\mathbf{A}-\mathbf{C})\cdot(\mathbf{A}-\mathbf{C})-r^2=0$

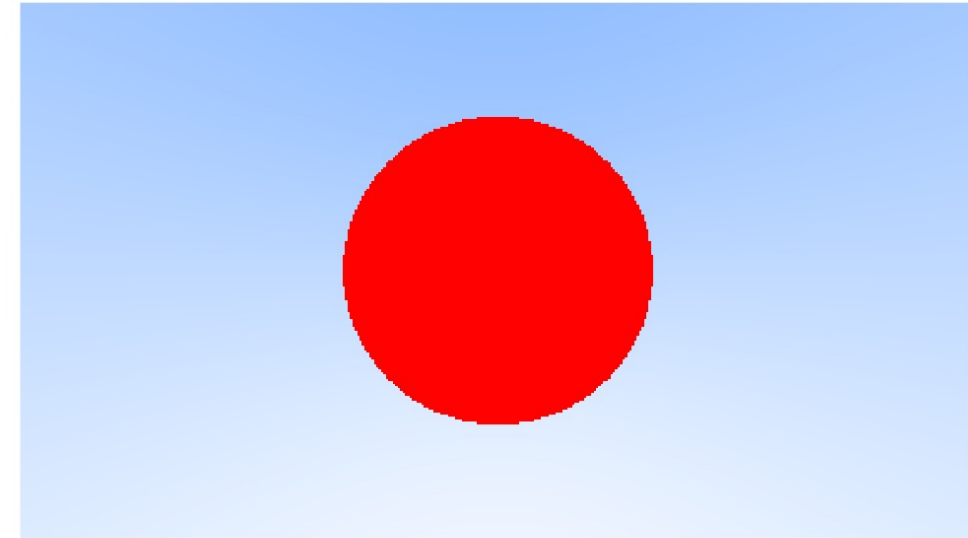


Image 3: A simple red sphere

```
bool hit_sphere(const point3& center, double radius, const ray& r) {  
    return (discriminant > 0);  
}
```

```
color ray_color(const ray& r) {  
    if (hit_sphere(point3(0,0,-1), 0.5, r))  
        return color(1, 0, 0);  
    vec3 unit_direction = unit_vector(r.direction());  
    auto t = 0.5*(unit_direction.y() + 1.0);  
    return (1.0-t)*color(1.0, 1.0, 1.0) + t*color(0.5, 0.7, 1.0);  
}
```

Surface Normals and Multiple Objects

- Surface normal: vector that is perpendicular to the surface at the point of intersection
- Code: configure `hit_sphere` to return smallest `t`.



Image 4: A sphere colored according to its normal vectors

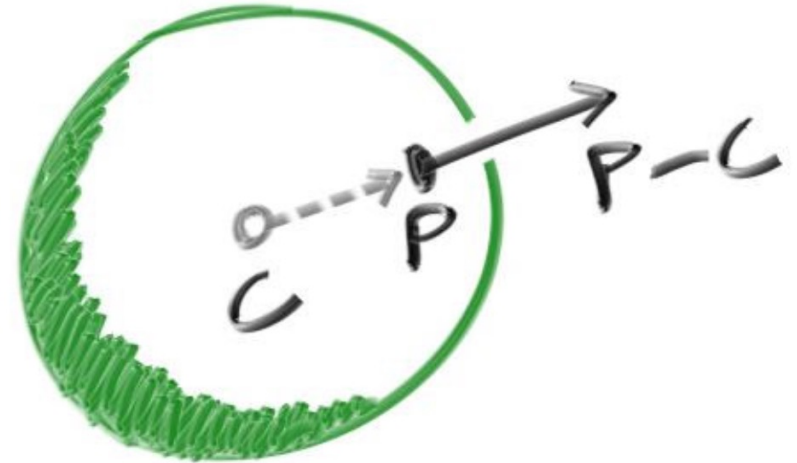


Figure 5: Sphere surface-normal geometry

An Abstraction for Hittable Objects

- Object, surface, hittable
- Struct hit_record: P, N, t
- Normal directions
 - Always out
 - Against ray

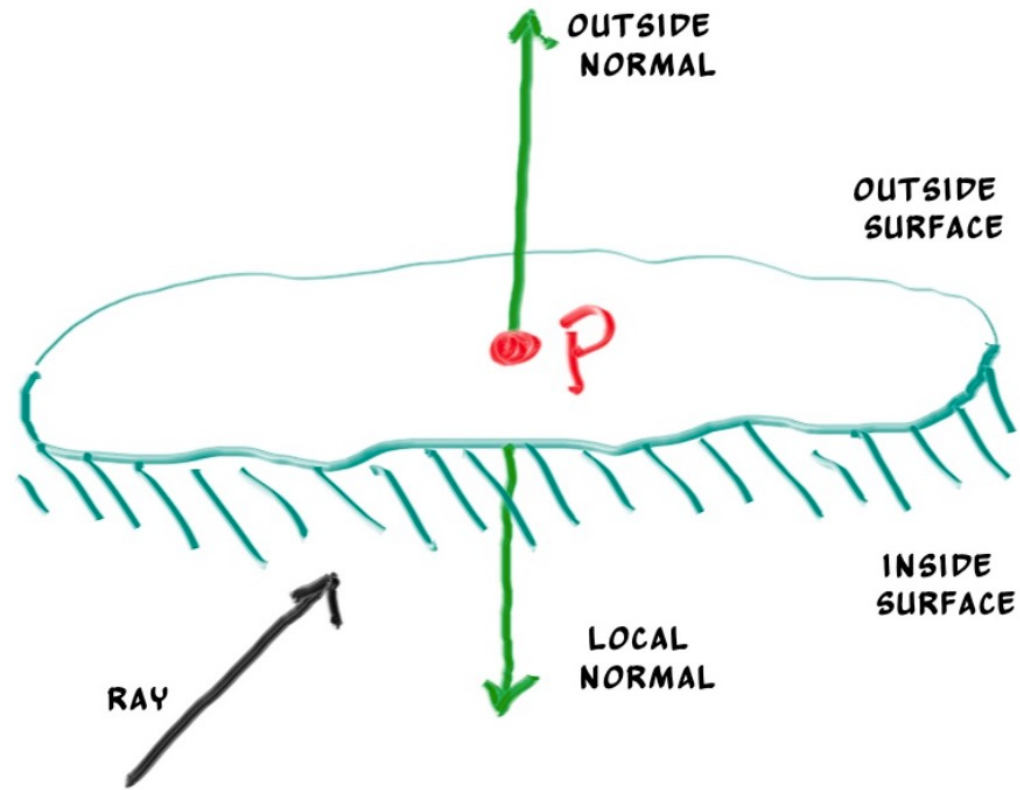


Figure 6: Possible directions for sphere surface-normal geometry

Add hittable_list

```
Vector3 ray_color(Ray r, hittable_list world) {  
    hit_record rec;  
    if (world.hit(r, 0, infinity, rec)) {  
        return 0.5 * (rec.N + Vector3(1,1,1));  
        std::cout<<"hit"<<std::endl;  
    }  
    Vector3 unit_direction = r.direction().unit();  
    auto t = 0.5*(unit_direction.y() + 1.0);  
    return (1.0-t)*Vector3(1.0, 1.0, 1.0) + t*Vector3(0.5, 0.7, 1.0);  
}
```

```
// World  
hittable_list world;  
world.add(make_shared<sphere>(point3(0,0,-1), 0.5));  
world.add(make_shared<sphere>(point3(0,-100.5,-1), 100));
```



Image 5: Resulting render of normals-colored sphere with ground

Antialiasing

- Real camera: no jaggies along edges because the edge pixels are a blend of some foreground and some background.
- => averaging a bunch of samples inside each pixel
- Canonical random number $0 \leq r < 1$
- => use `rand()` in `<cstdlib>` returns a random integer 0, RAND_MAX

```
✓ inline double random_double() {  
    // Returns a random real in [0,1).  
    return rand() / (RAND_MAX + 1.0);  
}  
  
✓ inline double random_double(double min, double max) {  
    // Returns a random real in [min,max).  
    return min + (max-min)*random_double();  
}
```

Add camera class

```
class camera {  
private:  
    Vector3 _origin;  
    Vector3 _lower_left_corner;  
    Vector3 _horizontal;  
    Vector3 _vertical;  
  
public:  
    camera();  
    Ray get_ray(double u, double v);  
};
```

```
inline double clamp(double x, double min, double max) {  
    if (x < min) return min;  
    if (x > max) return max;  
    return x;  
}
```

- Write `_color`: add the full color each iteration, and then perform a single divide at the end (by the number of samples) when writing out the color

Add samples

```
const int image_height = static_cast<ir  
const int samples_per_pixel = 100;
```

```
for (int i = 0; i < image_width; ++i) {  
    color pixel_color(0, 0, 0);  
    for (int s = 0; s < samples_per_pixel; ++s) {  
        auto u = (i + random_double()) / (image_width-1);  
        auto v = (j + random_double()) / (image_height-1);  
        ray r = cam.get_ray(u, v);  
        pixel_color += ray_color(r, world);  
    }  
    write_color(std::cout, pixel_color, samples_per_pixel);  
}
```

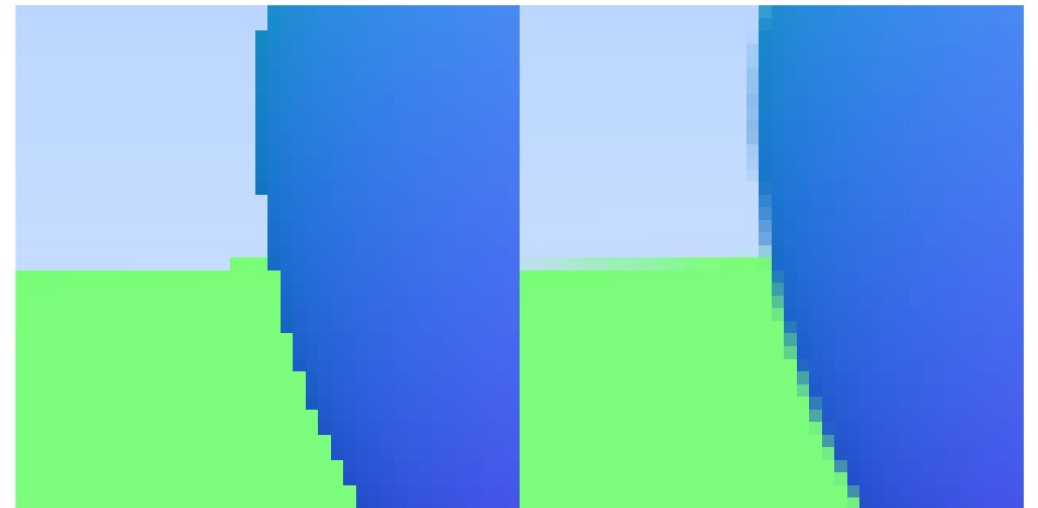


Image 6: Before and after antialiasing

Diffuse Materials

- take on the color of their surroundings, but they modulate that with their own intrinsic color
- Light that reflects off a diffuse surface has its direction randomized

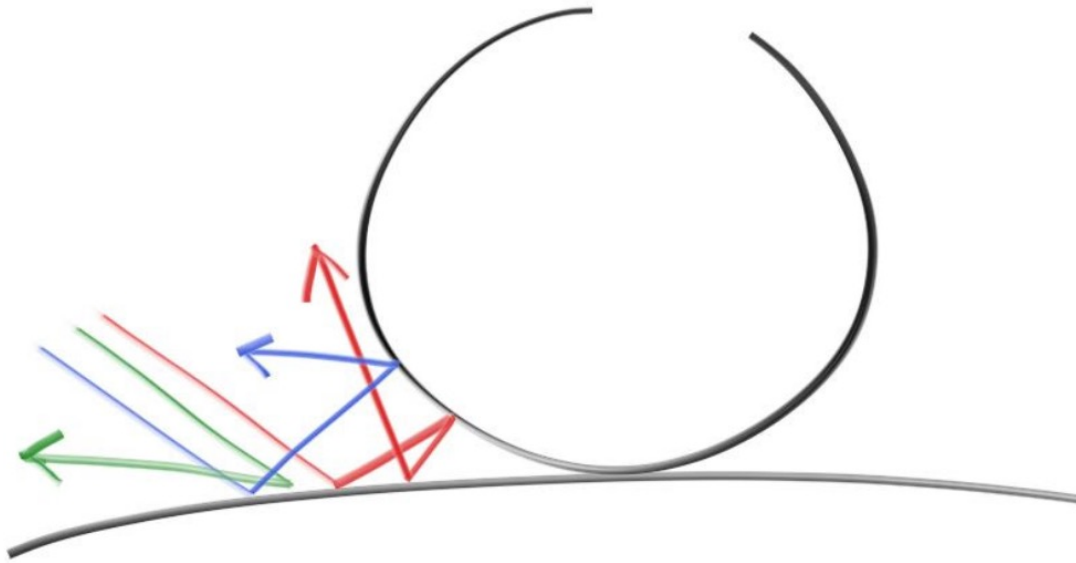
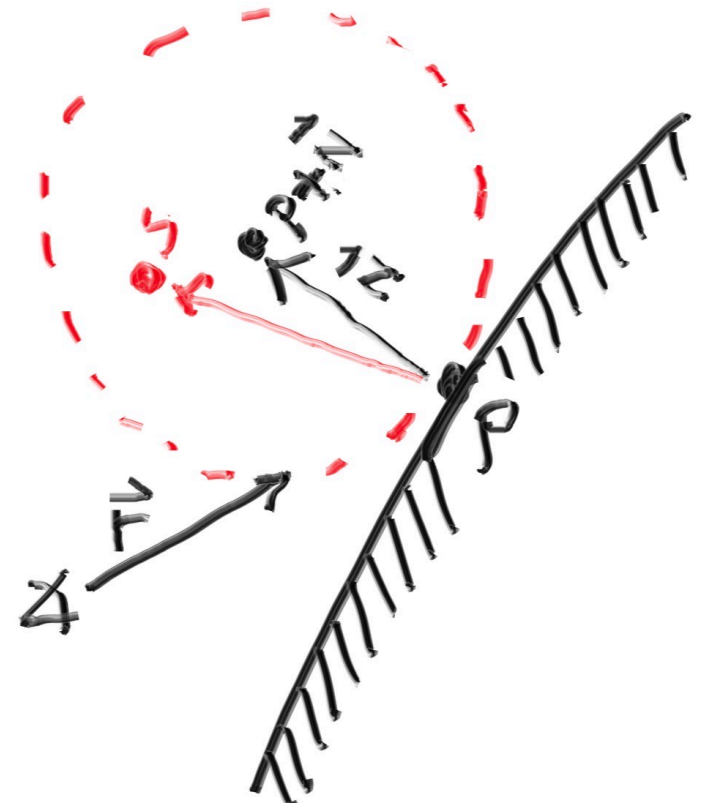


Figure 8: *Light ray bounces*

Diffuse Materials

- Also absorbed. The darker the surface, the more likely absorption is.
- algorithm randomizes direction will produce surfaces that look matte
- The sphere with a center at $(\mathbf{P}-\mathbf{n})$ is considered *inside* the surface, whereas the sphere with center $(\mathbf{P}+\mathbf{n})$ is considered *outside* the surface.
 - Select radius on the same side as the ray origin
 - Pick random \mathbf{S} inside this unit radius.



Pick random S.

- Rejection method:
- First, pick a random point in the unit cube where x, y, and z all range from -1 to +1. Reject this point and try again if the point is outside the sphere.

```
static Vector3 random();  
static Vector3 random(double min, double max);  
  
Vector3 random_in_unit_sphere();
```

Recursive ray_color

- Limit the number of child rays by depth.

```
Vector3 ray_color(Ray r, hittable_list world, int depth) {  
    hit_record rec;  
  
    if (depth <= 0) return Vector3(0, 0, 0);  
  
    if (world.hit(r, 0, infinity, rec)) {  
        Vector3 target = rec.P + rec.N + Vector3::random_in_unit_sphere();  
        return 0.5 * ray_color(Ray(rec.P, target - rec.P), world, depth - 1);  
    }  
}
```

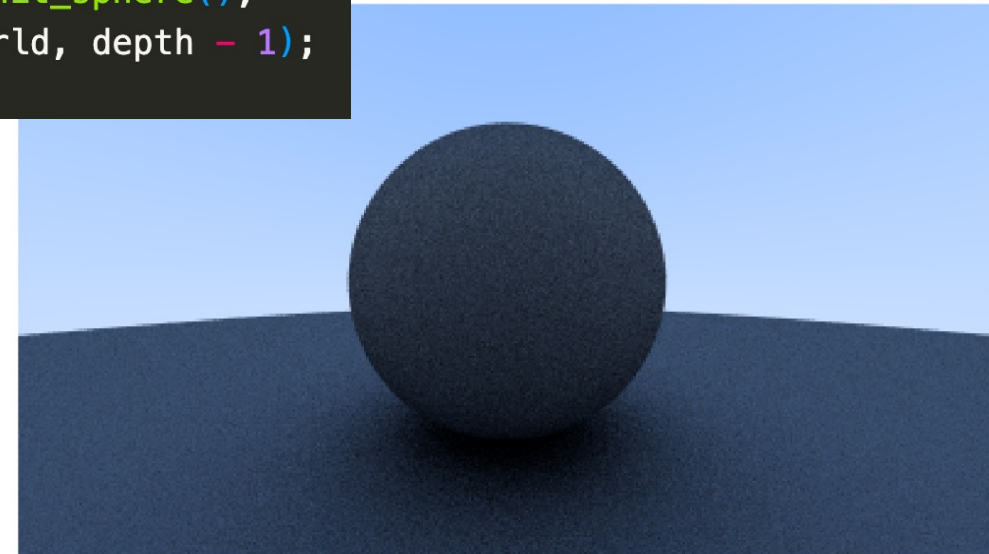


Image 7: First render of a diffuse sphere

Gamma correction for accurate color intensity

- meaning the 0 to 1 values have some transform before being stored as a byte
- use “gamma 2” which means raising the color to the power $1/\text{gamma}$, or in our simple case $\frac{1}{2}$, which is just square-root:

```
void write_color(std::ostream &out, color pixel_color, int samples_per_pixel) {  
    auto r = pixel_color.x();  
    auto g = pixel_color.y();  
    auto b = pixel_color.z();  
  
    // Divide the color by the number of samples and gamma-correct for gamma=2.0.  
    auto scale = 1.0 / samples_per_pixel;  
    r = sqrt(scale * r);  
    g = sqrt(scale * g);  
    b = sqrt(scale * b);  
  
    // Write the translated [0,255] value of each color component.  
    out << static_cast<int>(256 * clamp(r, 0.0, 0.999)) << ' '  
        << static_cast<int>(256 * clamp(g, 0.0, 0.999)) << ' '  
        << static_cast<int>(256 * clamp(b, 0.0, 0.999)) << '\n';  
}
```

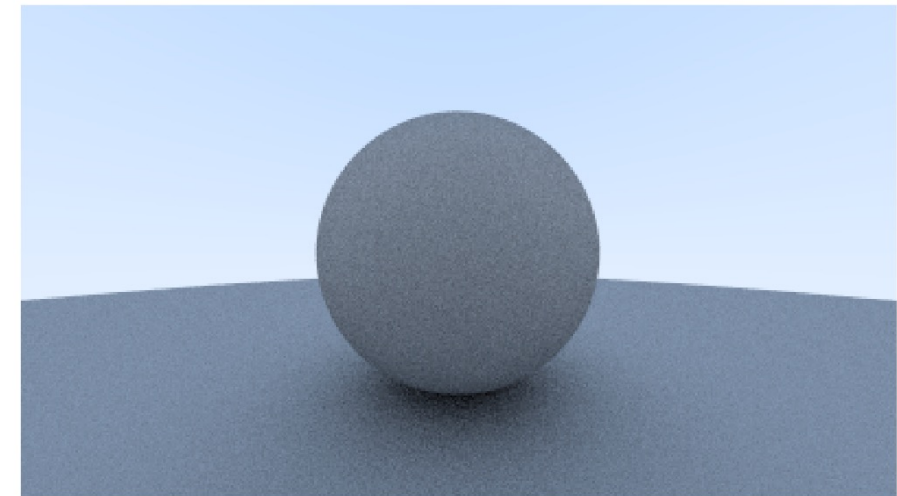
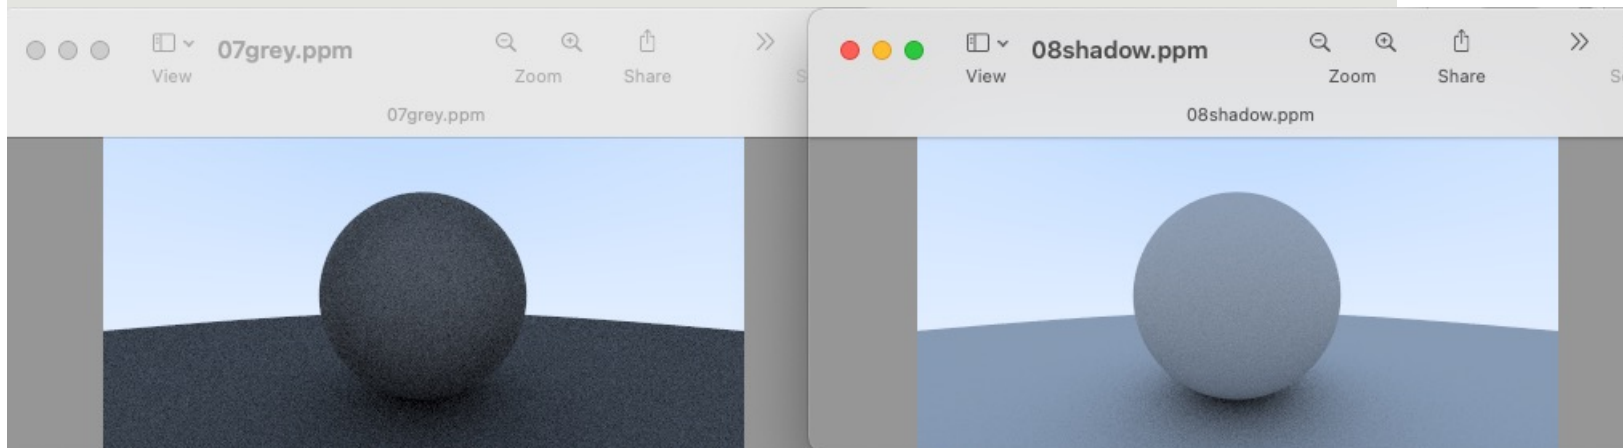


Image 8: Diffuse sphere, with gamma correction

Shadow acne problem

- Some of the reflected rays hit the object they are reflecting off of not at exactly $t=0$, but instead at $t=-0.00000001$ or $t=0.00000001$ or whatever floating point approximation the sphere intersector gives us. So we need to ignore hits very near zero:

```
if (world.hit(r, 0.001, infinity, rec)) {
```



True Lambertian Reflection

- Rejection method: scales by the $\cos^3(\phi)$ where ϕ is the angle from the normal. This is useful since light arriving at shallow angles spreads over a larger area, and thus has a lower contribution to the final color.
- Lambertian distribution: distribution of $\cos(\phi)$.
has the probability higher for ray scattering close to the normal, the distribution is more uniform.
picking random on the surface of the unit sphere,
offset along the surface normal.

```
Vector3 Vector3::random_unit_sphere() {  
    // on sphere S  
    return Vector3::random_in_unit_sphere().unit();  
}
```

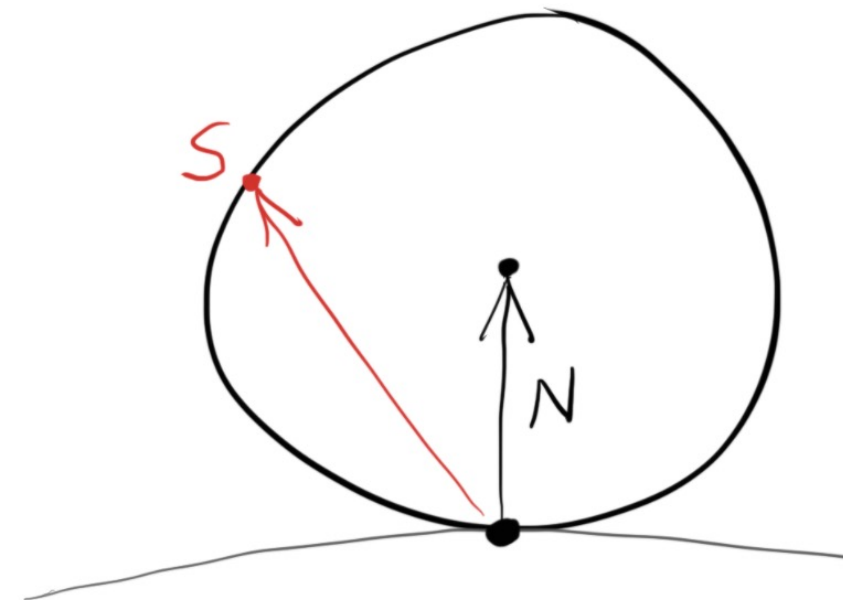
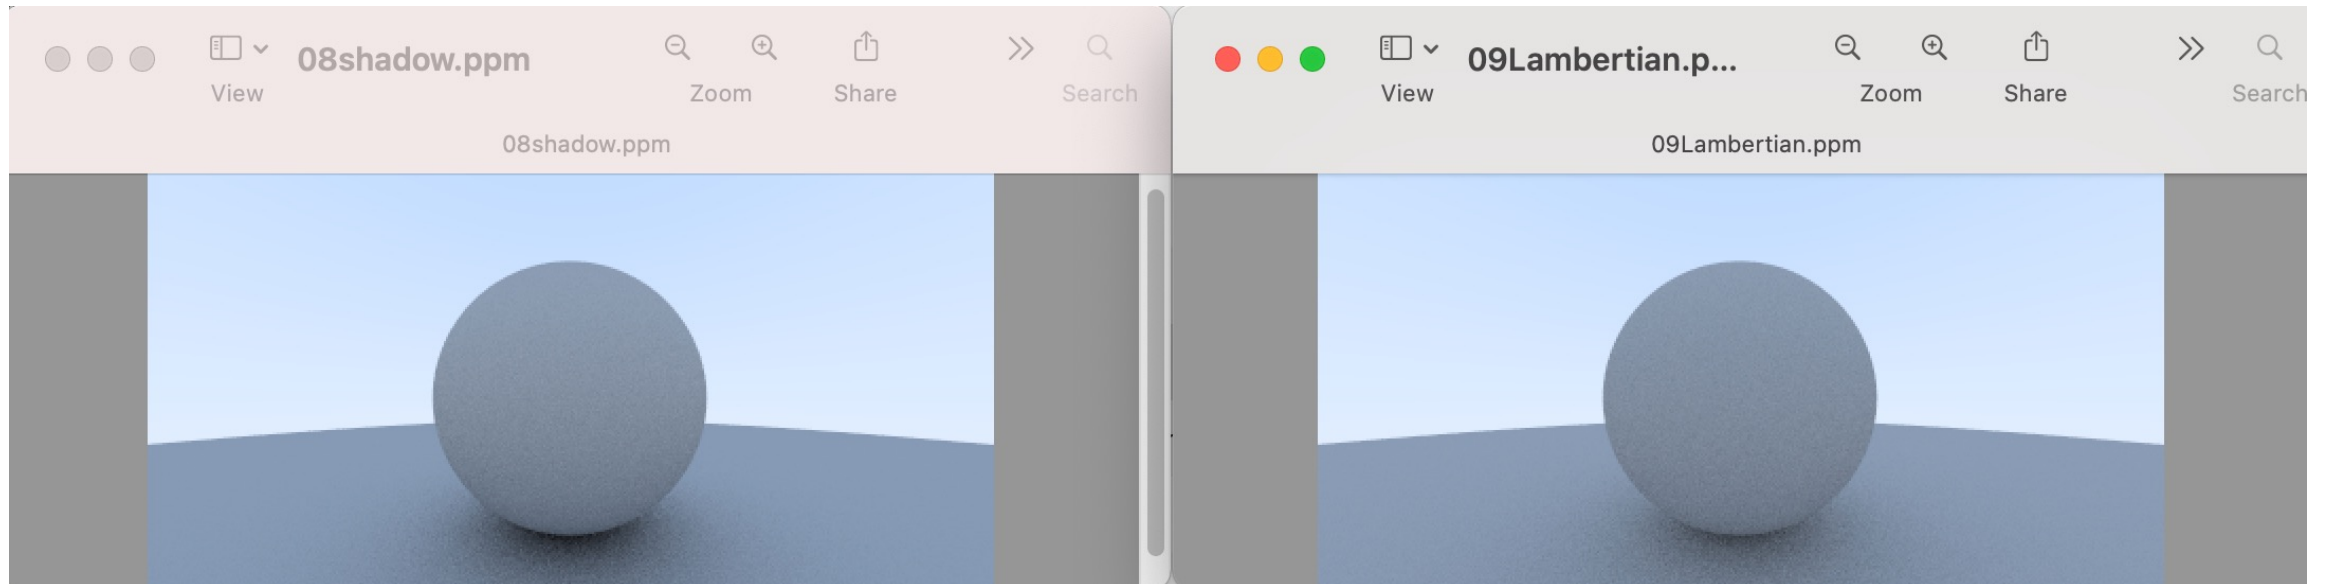


Figure 10: Generating a random unit vector

Differences

- The shadows are less pronounced after the change
 - Both spheres are lighter in appearance after the change
- <= due to the more uniform scattering of the light rays, fewer rays are scattering toward the normal, appear *lighter* because more light bounces toward the camera. the shadows less light bounces straight-up



Abstract class for materials

- Produce a scattered ray (or say it absorbed the incident ray).
- If scattered, say how much the ray should be attenuated.

```
#ifndef MATERIAL_H
#define MATERIAL_H

#include "rtweekend.h"

struct hit_record;

class material {
public:
    virtual bool scatter(
        const ray& r_in, const hit_record& rec, color& attenuation, ray& scattered
    ) const = 0;
};

#endif
```

Listing 41: [material.h] *The material class*

Material and hittable

- Hittables and materials need to know each other so there is some circularity of the references. In C++ you just need to alert the compiler that the pointer is to a class, which the “class material” in the hittable class below does:

```
#include "rtweekend.h"

class material;

struct hit_record {
    point3 p;
    vec3 normal;
    shared_ptr<material> mat_ptr;
    double t;
    bool front_face;
}
```

- the material pointer in the hit_record will be set to point at the material pointer, when hits.
- Add `shared_ptr<material> mat_ptr` to Sphere.

[hittable.h] Hit record with added material pointer

Modeling Light Scatter and Reflectance

scatter with some probability p and have attenuation be $albedo/p$.

```
class lambertian : public material {
public:
    lambertian(const color& a) : albedo(a) {}

    virtual bool scatter(
        const ray& r_in, const hit_record& rec, color& attenuation, ray& scattered
    ) const override {
        auto scatter_direction = rec.normal + random_unit_vector();

        // Catch degenerate scatter direction
        if (scatter_direction.near_zero())
            scatter_direction = rec.normal;

        scattered = ray(rec.p, scatter_direction);
        attenuation = albedo;
        return true;
    }

public:
    color albedo;
};
```

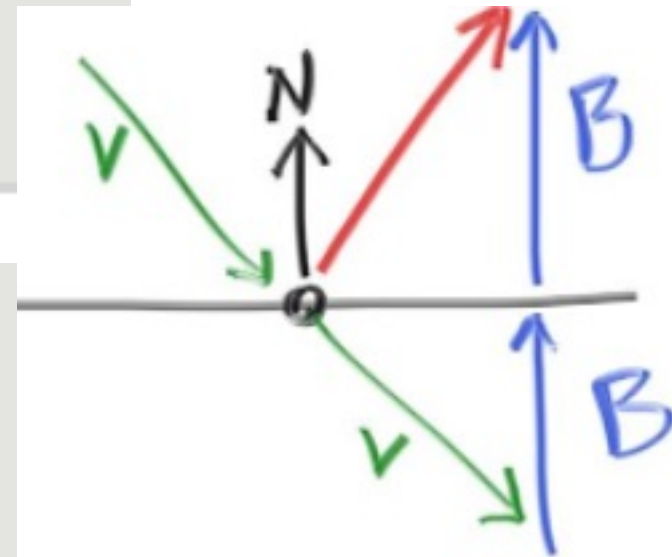
Listing 46: [material.h] *Lambertian scatter, bullet-proof*

Mirrored Light Reflection

```
vec3 reflect(const vec3& v, const vec3& n) {  
    return v - 2*dot(v,n)*n;  
}
```

Listing 47: [vec3.h] *vec3 reflection function*

```
class metal : public material {  
public:  
    metal(const color& a) : albedo(a) {}  
  
    virtual bool scatter(  
        const ray& r_in, const hit_record& rec, color& attenuation, ray& scattered  
    ) const override {  
        vec3 reflected = reflect(unit_vector(r_in.direction()), rec.normal);  
        scattered = ray(rec.p, reflected);  
        attenuation = albedo;  
        return (dot(scattered.direction(), rec.normal) > 0);  
    }  
  
public:  
    color albedo;  
};
```



Fuzz Mirror

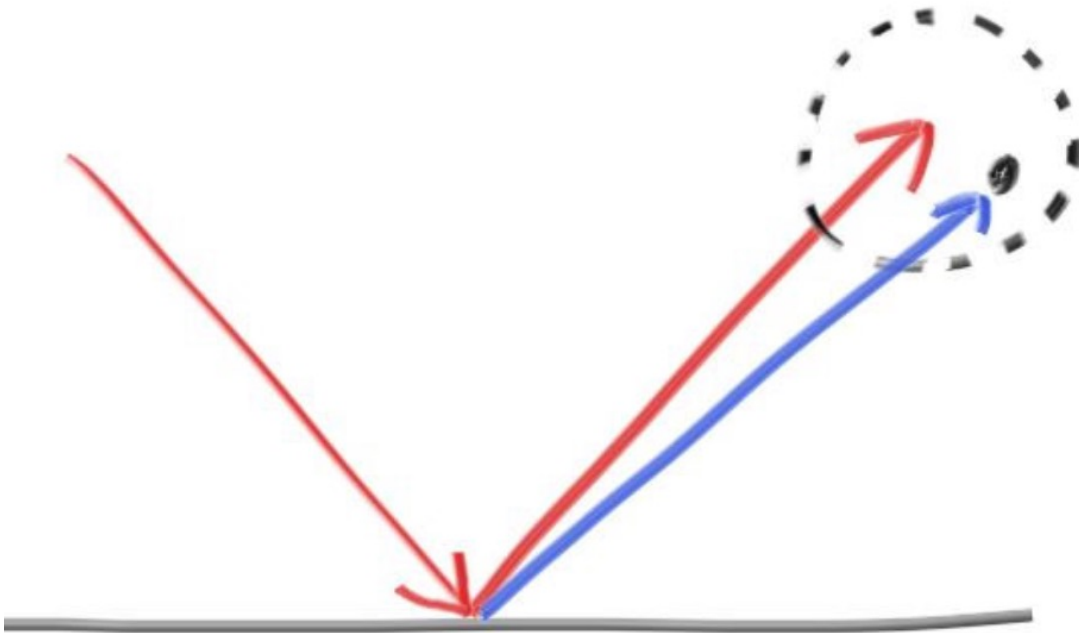
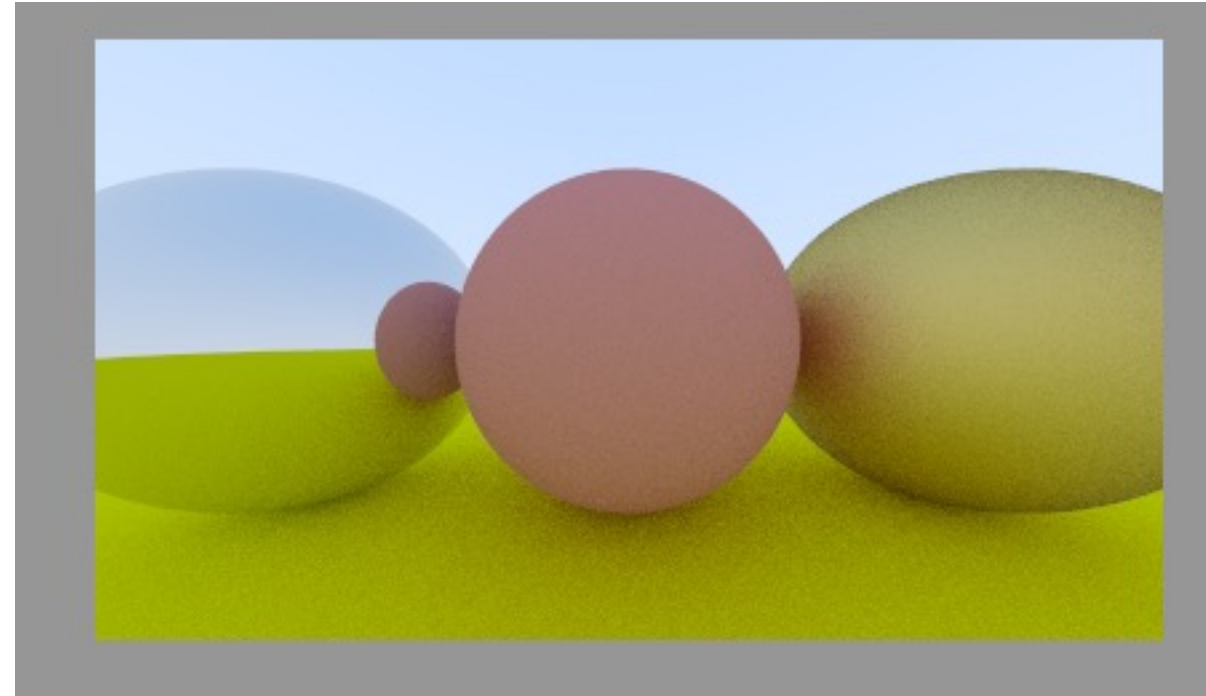


Figure 12: Generating fuzzed reflection rays



```
) const override {  
    vec3 reflected = reflect(unit_vector(r_in.direction()), rec.normal);  
    scattered = ray(rec.p, reflected + fuzz*random_in_unit_sphere());  
    attenuation = albedo;  
    return (dot(scattered.direction(), rec.normal) > 0);  
}
```

Fuzz = 0 & 1.